

COMP4906  
Honours Thesis Proposal  
A Real-Time Fluid Simulation for Fluid-Geometry  
Interaction in Game Environments

Gabriel Racz

November 15, 2024

# 1 Fluid Dynamics Background

## 1.1 Equations of Motion

A fluid with constant density and temperature can be modelled by a velocity vector field  $\mathbf{u}$  and a scalar pressure field  $p$ . These fields respectively represent the instantaneous velocity vector and outwards pressure force at any point within the fluid domain. A set of famous partial differential equations called the *incompressible Navier-Stokes* equations describe the evolution of an incompressible fluid over time.

$$\frac{\partial \mathbf{u}}{\partial t} + (\mathbf{u} \cdot \nabla) \mathbf{u} = -\frac{1}{\rho} \nabla p + \nu \nabla^2 \mathbf{u} + \mathbf{f} \quad (1)$$

$$\nabla \cdot \mathbf{u} = 0 \quad (2)$$

Equation (1) is analogous to Newton's momentum equation  $F = ma$  from classical mechanics. It describes the relationship between force and motion for a continuous fluid medium rather than discrete point particles. Equation (2) is the incompressibility constraint on the velocity field. It states that the *divergence* of the velocity field should be 0 at every point in space, ensuring that the fluid conserves mass over time.

Each term from the momentum equation governs a different aspect of fluid behavior which when combined produces complex flow patterns. Seeing as these are the basis for all preceding research, it is worth describing their individual contributions.

The simplest of the terms  $\mathbf{f}$  is the *external force*. This force incorporates outside influences acting on the system. Oftentimes this includes gravity, as well as artist-driven inputs.

The *viscous* term  $\nu \nabla^2 \mathbf{u}$  accounts for the internal friction within the fluid. It acts to diffuse the fluid velocity overtime, causing neighbouring regions to “drag” against each other, gradually bringing the system to a more uniform state. The Laplacian operator  $\nabla^2$  when applied to  $\mathbf{u}$  yields a new vector field representing the spatial variation of the velocity. Multiplying by the fluid's viscosity constant  $\nu$  updates local regions to be more inline with their neighbours.

$-\frac{1}{\rho} \nabla p$  is the pressure gradient term. It describes the change in fluid velocity due to the differences in pressure inside of the fluid. The pressure gradient  $\nabla p$ , force per unit volume, is divided by the fluid density  $\rho$  to obtain the acceleration due to this difference in pressure (similar to  $a = \frac{F}{m}$ ). The term is negated as the fluid naturally moves from areas of high to low pressure. As we will show later, this term is crucial in satisfying the incompressibility constraint (2)

Finally the convective term  $(\mathbf{u} \cdot \nabla) \mathbf{u}$  represents the transport of momentum due to the fluid's own motion. Commonly called *advection*, this models

the movement of the fluid along the velocity field. The differential operator  $(\mathbf{u} \cdot \nabla)$  encodes how the velocity changes as you follow a fluid particle's path, which we then apply to the velocity field itself. Notably this results in a nonlinear as the evolution of the velocity field overtime depends on its own spatial derivatives.

## 1.2 Incompressibility

In computer graphics applications, fluid simulations are generally performed under the assumption of incompressibility, where the density of the fluid remains constant throughout the flow. This vastly simplifies the Navier-Stokes equations for general compressible fluids which are otherwise complex and expensive to compute, yielding the simplified form described above. Under non-extreme circumstances, the effects of compressibility are negligible for the simulation of visually appealing flow behaviors. To ensure that the fluid satisfies the incompressibility constraint (2) however, we must ensure that updates made to the system via the momentum equation (1) also maintains the divergence-free property. This step is commonly referred to as the projection step. The intermediate velocity field  $\mathbf{u}^*$  obtained from solving the previously discussed advection, viscous, and external terms is *projected* onto a divergence-free representation of itself.

Various methods have been proposed for the solution of this projection step with various tradeoffs in physical accuracy and computational efficiency. One early method leveraged by Jos Stam [22] relies on the Helmholtz-Hodge Decomposition theorem, stating that any vector field  $\mathbf{w}$  can be uniquely decomposed into the form

$$\mathbf{w} = \mathbf{u} + \nabla q \quad (3)$$

where  $\mathbf{u}$  has zero divergence and  $q$  is a scalar field, effectively splitting the field  $\mathbf{w}$  into its *curl-free* and *gradient-free* components. The equation can be rewritten in two terms by multiplying both sides by  $\nabla$ , removing the  $\mathbf{u}$  term since  $\nabla \cdot \mathbf{u} = 0$ .

$$\nabla \cdot \mathbf{w} = \nabla^2 q \quad (4)$$

This equation is a Poisson Equation for the field  $q$ , the solution of which allows us to “correct” our intermediate velocity field  $\mathbf{u}^*$  by removing its divergence component. Stam used this strategy to define a projection operator  $\mathbf{P}$  that computes the projection  $\mathbf{u}$ :

$$\mathbf{u} = \mathbf{P}\mathbf{w} = \mathbf{w} - \nabla q$$

Applying the operator  $\mathbf{P}$  to the momentum equation (1), and using the fact that  $\mathbf{P}\mathbf{u} = \mathbf{u}$  due to the incompressibility constraint and  $\mathbf{P}\nabla p = 0$  as it is a pure gradient field, obtains a single equation governing the change in the

fluid velocity over time

$$\begin{aligned}\frac{\partial \mathbf{u}}{\partial t} &= \mathbf{P} \left( -(\mathbf{u} \cdot \nabla) \mathbf{u} + \nu \nabla^2 \mathbf{u} + \mathbf{f} \right) \\ \frac{\partial \mathbf{u}}{\partial t} &= \mathbf{P} (\mathbf{u}^*) - \nabla q\end{aligned}\tag{5}$$

Later, authors such as Fedkiw [5] and Bridson [1] applied a similar method to obtain projection operators that repurposed the pressure term from the momentum equation to analogously replace  $q$  in Stam's derivation. Namely, Fedkiw incorporated the change in time  $\Delta t$  with:

$$\mathbf{u} = \mathbf{u}^* - \Delta t \nabla p$$

and Bridson incorporated the fluid density

$$\mathbf{u} = \mathbf{u}^* - \frac{\Delta t}{\rho} \nabla p$$

Solving the Poisson equation of the form (4) to obtain the corrective pressure gradient is largely the most expensive step for a typical fluid solve [1], requiring efficient parallelization and discretization of the fluid domain.

### 1.3 Lagrangian and Eulerian Viewpoints

With the continuum of the fluid defined above, tracking its motion and evolution over time requires a discretization of the domain. For fluids, there are two categories that each lead to different methods for solving the Navier-Stokes equations: the Lagrangian frame and the Eulerian frame.

The Lagrangian viewpoint treats the fluid volume as a particle system where each point in the fluid is represented as a separate particle with position  $\vec{x}$  and velocity  $\vec{v}$ . The fluid behaviour is then governed by the interaction of these discrete particles. In computer graphics, the particle system updates the particles similar to classical mechanics, where the collisions and other local forces between molecules are integrated over time to update the system.

The Eulerian viewpoint instead looks at fixed points in space and measures how the fluid quantities such as density, and velocity change over time. This can be thought of as discretizing the fluid domain into a static grid, where each grid cell contains measurements on what the underlying fluid is doing at a given time.

Both of these methods give rise to different strategies when solving the fluid equations and are each suited to different applications. The Lagrangian approach naturally lends itself to the simulation of liquid, as the particle-based representation makes it straightforward to capture the free surface of the fluid, individual droplets, and splashes. This is mainly due to the

minimal computational overhead required to track the fluid interface when compared with a grid based approach, as well as liquid particles behaving more closely to a solid at the microscopic level. Methods such as SPH [14] and vortex methods [18] smooth these discrete particles together to form a continuous medium, while the updates are performed at a per-particle basis. The Eulerian viewpoint are often more commonly used in computer graphics due to the easier computation of local quantities such as the pressure gradient and diffusion. It is significantly easier to take spatial derivatives on a fixed grid than in a cloud of arbitrarily moving particles. Many methods leverage the strengths of both by using a hybrid approach such as PIC/FLIP[21] where particles carry fluid quantities but are transferred to an underlying grid for the computation of spatial derivations. In modelling incompressible flow, a semi-Lagrangian approach is taken to perform stable advection of the fluid, simplifying the computation of the advection term.

## 2 Literature

### 2.1 Stable Incompressible Fluids

Stam’s Stable Fluids paper [23] builds upon the work of Foster and Metaxas [7] to provide the first stable numerical solution to the incompressible Navier-Stokes equations for the purposes of real-time applications. His insight was to use the *method of characteristics* to solve the self-advection of the fluid by adopting a semi-Lagrangian frame. Rather than solving the nonlinear advection term using finite differencing as was done by Foster and Metaxas, he traced the path of representative particles, centered on each grid cell, to compute the motion of the fluid due to its own flow. Tracing grid particles forwards in time however leads to numerical instability due to accumulating interpolated velocities at the particles’ cell destinations. Instead, he proposed to trace the particle’s motion backwards in time along the velocity field to find the representative particle that would end up at each grid cell’s center. This has the benefit of ensuring each grid cell is updated by exactly one particle which is perfectly aligned to the fluid grid, and moving the interpolation to the representative particle starting location step. The advection term is updated by:

$$\mathbf{u}_{\text{new}} = \mathbf{u}(\mathbf{p}(\mathbf{x}, -\Delta t))$$

where  $\mathbf{p}$  returns the position of the the particle at  $\mathbf{x}$  advected backwards in time through the velocity field  $\mathbf{u}$ . He claims that this improvement makes the self-advection unconditionally stable for all velocities and time steps, and extremely desirable property for real-time fluids. Almost all preceding research on fluids for computer graphics leverages this powerful property.

Stam also introduces a method for visualizing the otherwise constant density fluid by incorporating a second dye, or smoke fluid quantity into the

system. This smoke is allowed to have variable density at different points in space, advecting along the velocity field  $\mathbf{u}$  with the same advection principal described above, as well as diffusing throughout the medium using the same viscous term found in (1)

In Visual Simulation of Smoke [5], Fedkiw et. al proposed various improvements on Stam’s method. Due to the nature of using a discretized grid to model a continuum such as a fluid, small-scale turbulent structures and high frequency details are lost due to numerical dissipation. Once details fall below a resolution that can “escape” a single grid cell, they are lost to the diffusive properties of the grid updates. To reintroduce these details to the fluid, Fedkiw et al. applied a method invented by Steinhoff [24] for CFD called *Vorticity Confinement*. In incompressible flow, the vorticity of the velocity field is defined by:

$$\boldsymbol{\omega} = \nabla \times \mathbf{u}$$

This is also called the curl of the velocity field  $\mathbf{u}$ . This vector field describes the local spinning motion of each region within  $\mathbf{u}$ , similar to a paddle wheel spinning the flow in a certain direction. We can leverage the curl to adding forces back into the system after that have been lost due to the damping in our coarse grid solutions, recovering some of the small-scale turbulent behavior. By taking the normalized gradient of the curl

$$\begin{aligned} \boldsymbol{\eta} &= \nabla |\boldsymbol{\omega}| \\ \mathbf{N} &= \frac{\boldsymbol{\eta}}{|\boldsymbol{\eta}|} \end{aligned} \tag{6}$$

We obtain the vectors that point from areas of low vorticity to areas of higher vorticity. By projecting back into the velocity field, we can add an amount of vorticity that would be transported by the field

$$\mathbf{f}_{\text{vorticity confinement}} = \epsilon h (\mathbf{N} \times \boldsymbol{\omega}) \tag{7}$$

where  $\epsilon$  is a user-controlled parameter indicating the desired strength of the confinement, and  $h$  is the grid cell spacing.

Fedkiw et.al also improve upon the stable advection interpolation as described by Stam by using a higher order cubic *Hermite* interpolation scheme rather than the simple linear interpolation implemented by Stam. By considering more control points when interpolating the particle’s initial position the numerical diffusion, (blurring) is reduced between solver time steps, further contributing to the preservation of small-scale turbulence. Other interpolation schemes such as B-Spline [11] can be substituted here, further increasing accuracy while sacrificing computational efficiency.

He et. al built upon this work by introducing a method for spatially adaptive vorticity confinement, where the vorticity is used to up-sample a coarse velocity grid to improve performance. They also augment the vorticity confinement to be proportional to the *helicity* of the fluid volume  $V$

defined as:

$$H = \int \mathbf{u} \cdot \boldsymbol{\omega} dV$$

which is a scalar field that represents the twisting of the flow around the velocity vectors. This serves to add a spatially varying confinement coefficient to the previously defined vorticity confinement equation when discretized

$$\mathbf{f}_{\text{vorticity confinement}} = \epsilon h |\mathbf{u} \cdot \boldsymbol{\omega}| \left( \mathbf{N} \times \frac{\boldsymbol{\omega}}{|\boldsymbol{\omega}|} \right) \quad (8)$$

They use this more accurate vorticity confinement method to allow discretizing the velocity field over a relatively coarser grid by up sampling it and retrieving the lost turbulent forces due to the increased damping caused by the lower resolution.

As shown previously show, solving for the pressure force required to maintain incompressibility between velocity field updates requires solving the poisson equation

$$\nabla^2 p = \nabla \cdot \mathbf{u}^*$$

When discretized to a grid, this equation is equivalent to a linear system

$$Ap = b$$

where  $A$  is the Laplacian operator  $\nabla^2$ ,  $p$  is the unknown pressure field we are solving for, and  $b$  is the velocity divergence field  $\nabla \cdot \mathbf{u}$ . Bridson shows the derivation for the solution to this linear system being equivalent to a finite differencing in each spatial direction [1].

The divergence of the velocity field represents the net flow entering/exiting each cell of  $\mathbf{u}$ . Stam implemented simple central finite differencing in his 2D fluid simulation as using the following formula:

$$\nabla \cdot \mathbf{u}(x, y) = \frac{\mathbf{u}(x + h, y) - \mathbf{u}(x - h, y)}{2h} + \frac{\mathbf{u}(x, y + h) - \mathbf{u}(x, y - h)}{2h}$$

where  $h$  is the grid cell separation. While efficient and simple to implement, Bridson notes that this formulation has a a glaring numerical weakness [1] – the value of  $\mathbf{u}(x, y)$  is never considered when computing its own divergence. This makes it possible for the cell  $\mathbf{u}(x, y)$  to be quite different from its neighbours while still seemingly reporting a derivative of 0. Seeing as though it is quite possible for the fluid domain to exhibit sharp variations such as these, we need a method that will report the divergence more accurately to enforce incompressibility. Ideally, we would like to be able to compute

$$\nabla \cdot \mathbf{u}(x, y) = \frac{\mathbf{u}(x + h/2, y) - \mathbf{u}(x - h/2, y)}{h} + \frac{\mathbf{u}(x, y + h/2) - \mathbf{u}(x, y - h/2)}{h}$$

however this would lead to sampling the velocity grid along its separating faces, and not in the cell centers where the velocity is stored. To obtain an accurate and unbiased second-order central difference, Bridson employs the use of a MAC grid [10].

This grid structure staggers the velocity grid relative to the pressure grid, storing the velocity vector components on the corresponding faces of their pressure grid cell.




Figure 1: One cell from the three-dimension MAC grid as published in [1]

The MAC grid structure greatly improves the stability of the pressure calculation and subsequent incompressibility constraints, however it requires interpolation to evaluate the full velocity vector at any point. Despite necessary interpolation, most modern incompressible fluid solvers, real-time or otherwise use this representation.

The Laplacian matrix  $A$  is a sparse symmetric which encodes the computation of the partial derivatives for each of the pressure cells within  $p$  using a similar finite differencing scheme. Each row  $i$  of the Laplacian computes the partial derivative of its corresponding pressure cell. In 3D, the diagonal of the Laplacian stores coefficient  $-6$  (as each cell has 6 neighbours to difference), and each row contains mostly 0s, except in the columns corresponding to that row's neighbours where 1s are stored. Seeing as it is so sparse, it is prohibitively inefficient to store the entire Laplacian in memory – in three dimensions its size is cubic relative to the size of the pressure matrix – the solution of the Poisson equation is performed iteratively.

The solution to systems containing five- or seven-point Laplacian problems have been exhaustively studied, as they appear in countless problems throughout engineering and science. Thus, there are many numerical methods available, each with tradeoffs to performance and accuracy. Simple methods such as various forms of the Jacobi iteration method [8] Gauss-Seidel relaxation [8], and Successive Over-Relaxation [9], have often been used in computer graphics applications where accuracy is not a priority applications due to their low complexity, however they suffer in convergence



time to an optimal solution. Other popular methods for modern fluid simulation in VFX include the Conjugate Gradient Method, Multigrid, and Domain Expansion [1]. This area requires further experimentation towards methods that best satisfy the performance and accuracy requirements for real-time fluids.

Once the pressure field  $p$  is found, we subtract it from the intermediate velocity field  $\mathbf{u}^*$  to obtain the final incompressible fluid result for the timestep.

## 2.2 Boundary Conditions

Handling interaction between solid geometry and fluids requires satisfying additional constraints in the cells that contain solid structures. To prevent fluid flow from entering a solid voxel cell we must ensure that the component of the fluid velocity normal to the solid voxel boundary must be 0

$$\vec{u} \cdot \vec{n} = 0$$

If the solid is moving through the fluid, we can update the constraint to allow the fluid velocity to match the normal component of the solid’s velocity.

$$\vec{u} \cdot \vec{n} = \vec{u}_{\text{solid}} \cdot \vec{n}$$

Since we are only limiting the fluid’s velocity component relative to the normal of the solid boundary, this is called the *no-stick* condition. It says the tangential velocity of the fluid is unaffected by the solid boundary, and assumes inviscid flow around it. Since the solid cell contains no fluid, and the pressure field indicates how much force should be applied at each cell to force the system to be incompressible, we also must populate solid boundary cells with  $p = 0$  to ensure that it does not exert a corrective force on neighbouring fluid cells. Setting values of the boundaries directly is commonly known as a *Dirichlet* boundary condition. Although straightforward, Dirichlet boundary conditions have been shown to exhibit artifacts in the resulting fluid solutions such as artificial pressure gradients at the boundaries. Modern practice is to enforce *Neumann* boundary conditions on the pressure:

$$\frac{\partial p}{\partial \vec{n}} = \nabla p \cdot \vec{n}$$

which specifies the normal derivative of the pressure rather than the direct value. To handle this, some have suggested applying fluid quantities in areas that do not contain actual fluid (such as solid boundaries) [1, 6]. These are usually referred to as *ghost cells*, and allow us to use the solvers of the continuum equations to implicitly handle other non-fluid behaviors. Bridson describes manually setting the velocity and pressure values on solid boundaries in such a way that the solution of the pressure gradient handles the

Neumann boundary conditions between the fluid and solid. His derivation reveals that setting the pressure at the solid boundaries to:

$$p_{\text{solid}} = p_{\text{fluid}} - \frac{\rho \Delta x}{\Delta t} (u_{\text{fluid}} - u_{\text{solid}})$$

creates a Neumann boundary condition that is handled by the same equations that govern the fluid interior, assuming that  $p_{\text{fluid}}$  is a neighbouring fluid cell to  $p_{\text{solid}}$  and  $u$  is the separating MAC grid velocity face between them at the same point in space. As previously mentioned,  $\frac{\rho \Delta x}{\Delta t}$  is the scaling term applied to the pressure Laplacian to account for density, timestep, and grid cell size.

This “voxelized” approach to handling solid-fluid interaction is efficient and suits real-time applications that have considerable amounts of dynamic geometry. With sufficiently tessellated meshes, labelling of the solid voxels within the fluid structure can be performed at the cost of a single point lookup on the geometry vertices. Otherwise, fast voxelization methods for rasterizing input geometry directly into a dynamic grid structure can be applied described in [12] or [4].

### 2.3 Sparse Dynamic Grids

Uniform grids are often used for most basic 3D solvers due to their extremely fast data access, contiguous memory layout, and ease of implementation. However, they also carry several limitations that make them impractical for real-time applications. One of these drawbacks is illustrated when simulating scenes where the region of fluid occupies a relatively small area of the total simulation domain. Large amounts of computation and memory are wasted on regions that have little to no bearing on the updates to the active fluid region. This is less of a problem for liquids, where a tight bounding box can be translated between updates by predicting the motion of the liquid volume through its previous velocity field. However, for simulation of gases, areas of seemingly empty or inactive space still significantly contributes to the solutions of pressure. To remedy this, various data structures to store *sparse* representations of the fluid grid have been developed.

The simplest approach as described by Bridson is to use a sparse blocked grid. In this implementation, he implemented a two-level grid hierarchy where a coarse grid was used to tile to simulation domain, and subsequent “fine blocks” – storing tiny contiguous arrays – could be assigned to areas of high detail. He later extended his work by increasing the tree depth to 4 levels, and adding a hashing index scheme to minimize the overhead introduced by the tree index traversal [3].

Losasso et al. [16] implemented an octree data structure to adaptively refine their simulation grid using various refinement criteria. They described

the following modifications to the Navier-Stokes equations in order to account for the variation in volume between grid cells

$$V_{\text{cell}} \nabla \cdot (\mathbf{u}^* - \Delta t \nabla p) = 0 \quad (9)$$

$$V_{\text{cell}} \nabla \cdot (\Delta t \nabla p) = \sum_{\text{faces}} ((\Delta t \nabla p)_{\text{face}} \cdot \mathbf{n}) A_{\text{face}} \quad (10)$$

Where  $V_{\text{cell}}$  is the volume of a of their largest cell, is used to construct a linear system to solve for pressure considering the area of each adjacent face. They refined their octree structure prioritizing areas containing solid geometry, vorticity at each grid cell, and a clamped density range just above the visible limit while being low enough to warrant higher resolution.

Similarly, Lentine et. al [15] applied a similar hierarchical tree structure with two main differences. They subdivided a coarse grid representation into finer uniform grids which allowed them to maintain efficient linear truncated access at each grid level. They also opted to solve the intermediary velocity updates on the finest grid resolution, and only utilized the hierarchical representation to compute the pressure projection locally at for each grid resolution. To supplement the detail lost, they also added Kolmogorov noise as done in [20] during upsampling of the fine grid. Various other noise functions have been used to improve down sampling detail from sparse grids and should be investigated [13, 2].

Modern approaches to fluid simulation largely rely on the GPU to accelerate highly parallel nature of the Eulerian Navier-Stokes equations. A family of methods focused on optimizing sparse dynamic grids in VRAM have been developed in recent years based on the original paper VDB [17]. DCGrid [19] is the newest of these and implements a sparse adaptive grid that incorporates the paging structure of modern GPU memory layouts. They put particular emphasis on creating mip-mapped resolutions of the grid structure, as well as utilizing apron cells around boundaries of changing resolutions to optimize stencil operations which are common when solving the pressure and diffusion equations.

## 2.4 Rendering

Rendering volumetric data requires a different approach to traditional geometry. The most common approach for rendering volumetric data is to use ray marching. In the case of gaseous fluids such as smoke and fire, rays marched along the fluid domain, accumulating radiance of each fluid cell they traverse.

Rasmussen et. al [20] introduce a few interesting improvements over standard ray marching. They modified their simulation domain to be cylindrical inline with the camera view frustum rather than rectangular, resulting in each grid cell being a truncated pyramid. This has two ingenious benefits:

the ray marching algorithm which usually needs to perform a line-rasterizing algorithm to traverse the grid cells is instead replaced by simple cell index increments, and the refinement of the simulation domain near the camera view is handled implicitly by the increasing cell sizes relative to view distance. Finally, they mention that for high albedo medium such as smoke, light scattering within the fluid volume becomes effectively isotropic and can be modeled as a diffusion process. This along with other fluid quantities such as temperature can be diffused by the efficient fluid solvers already in place.

### 3 Objective

The objective of my thesis will be to develop a framework for efficient GPU gaseous fluid simulation for the use in real-time games with a focus on geometry-fluid interaction. This will include the development of a dynamic player-centered sparse grid which will be incrementally refined around game geometry and areas of high fluid detail. Other requirements are as follows:

- Create visually interesting physically-based fluid flow around complex geometry
- End-to-end computational and memory efficiency suitable for integration into production game engines
- Volumetric rendering of fluid data incorporating light scattering
- GPU voxelization/simplification of scene geometry into the fluid domain
- Leverage modern GPU compute capabilities using Vulkan to implement methods found during literature review

## 4 Schedule

Period	Milestones
November 18 - December 1	Implement adaptive surface voxelization of triangle meshes (COMP5115 final project)
December 2 - December 15	Further literature review guided by proposal feedback
December 16 - December 29	Finalize geometry pipeline rendering and basic volumetric renderer
December 30 - January 12	Incorporate MAC grid finite differencing and Neumann boundary conditions from [1]
January 13 - January 26	Augment fluid simulation to 3D
January 27 - February 9	Update uniform grid structures and algorithms to support sparse dynamic grids
February 10 - February 23	Begin implementation for adaptive dynamic grid and refinement around fluid
February 24 - March 9	Finalize adaptive dynamic grid and improve volumetric renderer
March 10 - March 23	Performance optimization, compile results for report
March 24 - April 6	Write and submit final report draft
April 7 - April 20	Update final report based on feedback, submit

## References

- [1] Robert Bridson. *Fluid simulation for computer graphics*. AK Peters/CRC Press, 2015.
- [2] Robert Bridson, Jim Houriham, and Marcus Nordenstam. “Curl-noise for procedural fluid flow”. In: *ACM Transactions on Graphics (ToG)* 26.3 (2007), 46–es.
- [3] Robert Edward Bridson. *Computational aspects of dynamic surfaces*. stanford university, 2003.
- [4] Shiaofen Fang and Hongsheng Chen. “Hardware accelerated voxelization”. In: *Computers & Graphics* 24.3 (2000), pp. 433–442.
- [5] Ronald Fedkiw, Jos Stam, and Henrik Wann Jensen. “Visual simulation of smoke”. In: *Proceedings of the 28th annual conference on Computer graphics and interactive techniques*. 2001, pp. 15–22.
- [6] Ronald P Fedkiw et al. “A non-oscillatory Eulerian approach to interfaces in multimaterial flows (the ghost fluid method)”. In: *Journal of computational physics* 152.2 (1999), pp. 457–492.
- [7] Nick Foster and Dimitris Metaxas. “Modeling the motion of a hot, turbulent gas”. In: *Proceedings of the 24th annual conference on Computer graphics and interactive techniques*. 1997, pp. 181–188.
- [8] Gonalo Amador Abel Gomes. “Linear solvers for stable fluids: GPU vs CPU”. In: *17th EncontroPortugues de ComputacaoGrafica (EPCG09)* (2009), pp. 145–153.
- [9] A Hadjidimos. “Successive overrelaxation (SOR) and related methods”. In: *Journal of Computational and Applied Mathematics* 123.1-2 (2000), pp. 177–199.
- [10] Francis H Harlow, J Eddie Welch, et al. “Numerical calculation of time-dependent viscous incompressible flow of fluid with free surface”. In: *Physics of fluids* 8.12 (1965), p. 2182.
- [11] Shengfeng He et al. “Real-time smoke simulation with improved turbulence by spatial adaptive vorticity confinement”. In: *Computer Animation and Virtual Worlds* 22.2-3 (2011), pp. 107–114.
- [12] Christian Friedrich Janßen, Nils Koliha, and Thomas Rung. “A fast and rigorously parallel surface voxelization technique for GPU-accelerated CFD simulations”. In: *Communications in computational physics* 17.5 (2015), pp. 1246–1270.
- [13] Theodore Kim et al. “Wavelet turbulence for fluid simulation”. In: *ACM Transactions on Graphics (TOG)* 27.3 (2008), pp. 1–6.

- [14] Dan Koschier et al. “Smoothed particle hydrodynamics techniques for the physics based simulation of fluids and solids”. In: *arXiv preprint arXiv:2009.06944* (2020).
- [15] Michael Lentine, Wen Zheng, and Ronald Fedkiw. “A novel algorithm for incompressible flow using only a coarse grid projection”. In: *ACM Transactions on Graphics (TOG)* 29.4 (2010), pp. 1–9.
- [16] Frank Losasso, Frédéric Gibou, and Ron Fedkiw. “Simulating water and smoke with an octree data structure”. In: *Acm siggraph 2004 papers*. 2004, pp. 457–462.
- [17] Ken Museth. “VDB: High-resolution sparse volumes with dynamic topology”. In: *ACM transactions on graphics (TOG)* 32.3 (2013), pp. 1–22.
- [18] Sang Il Park and Myoung Jun Kim. “Vortex fluid for gaseous phenomena”. In: *Proceedings of the 2005 ACM SIGGRAPH/Eurographics symposium on Computer animation*. 2005, pp. 261–270.
- [19] Wouter Raateland et al. “Dcgrid: An adaptive grid structure for memory-constrained fluid simulation on the gpu”. In: *Proceedings of the ACM on Computer Graphics and Interactive Techniques* 5.1 (2022), pp. 1–14.
- [20] Nick Rasmussen et al. “Smoke simulation for large scale phenomena”. In: *ACM SIGGRAPH 2003 Papers*. 2003, pp. 703–707.
- [21] Fredrik Salomonsson. *PIC/FLIP Fluid Simulation Using Block-Optimized Grid Data Structure*. 2011.
- [22] Jos Stam. “Real-time fluid dynamics for games”. In: *Proceedings of the game developer conference*. Vol. 18. 11. 2003.
- [23] Jos Stam. “Stable fluids”. In: *Seminal Graphics Papers: Pushing the Boundaries, Volume 2*. 2023, pp. 779–786.
- [24] John Steinhoff and David Underhill. “Modification of the Euler equations for “vorticity confinement”: Application to the computation of interacting vortex rings”. In: *Physics of Fluids* 6.8 (1994), pp. 2738–2744.